

TECHNICAL LAB LOG

Session: 2026-07-09 to 2026-07-10

Author: Ramon E. Ramirez

Classification: Career Documentation

SESSION TIMELINE

Phase 1: Initial Assessment (22:00 - 22:30)

- Identified OpenHuman installation on E: drive
- Confirmed Ollama running with multiple custom models (ornith:35b, dhu-ornith-builder, llama3, gemma4, etc.)
- Verified GPU: NVIDIA RTX 4070 Ti detected via WMI
- Initial app launch: Window appeared but model selection unclear

Phase 2: Configuration Analysis (22:30 - 23:00)

- Located OpenHuman config: C:[user].openhuman[id].toml
- Analyzed TOML structure: 18KB configuration file
- Key findings:
 - default_model = “chat-v1” (cloud model)
 - local_ai.runtime_enabled = false (Ollama disabled)
 - onboarding_completed = false (setup wizard incomplete)
 - chat_onboarding_completed = false (chat blocked)
 - Cost tracking enabled with USD pricing tiers

Phase 3: First Config Edit (23:00 - 23:15)

- Attempted PowerShell-based config replacement
- Changed: default_model, chat_model_id, model_id to ornith:35b
- Set: onboarding_completed = true, chat_onboarding_completed = true
- Result: App launched but chat model blocked by MVP allowlist
- Log evidence: [local_ai] chat model not in MVP allowlist, redirecting to default resolved=“ornith:35b” fallback=“gemma3:1b-it-qat”

Phase 4: UI Verification (23:15 - 23:30)

- Confirmed app running via Task Manager (7 processes)
- Window state corrupted – 8 zombie processes detected
- Used taskkill /f /im OpenHuman.exe to clear processes
- Attempted window state reset via del window* commands
- Result: App opened but model showed ornith in UI while logs indicated cloud fallback

Phase 5: Governance Framework Introduction (23:30 - 00:00)

- User introduced Gen-Cultivate v1 framework (7 layers)
- Ingested full documentation: AAS-1, HPL, CVL, PCS, EMBR, GEN-CULTIVATE, CLEARANCE

- Applied framework to AI assistant reasoning:
 - HPL: Meaning Assertion before every suggestion
 - CVL: Self-verification of claims before amplification
 - PCS: Previous path consequences logged as experiential proof
 - EMBR: Pace regulation (user controls speed, not AI)
 - CLEARANCE: Explicit authorization required before execution
- Result: Measurable behavioral change in AI responses (precision increase, drift reduction)

Phase 6: Alternative System Evaluation (00:00 - 00:30)

- Evaluated AgenticSeek (GitHub: Fosowl/agenticSeek)
- Cloned to E:
- Analyzed Python architecture: FastAPI backend, per-agent prompts, Ollama integration
- Configured: config.ini -> provider_name = ollama, provider_model = ornith:35b
- Implemented governance injection at agent.py load_prompt() function
- Created governance.txt preamble prepended to all agent prompts
- Verified patch: load_prompt now loads governance before original prompt

Phase 7: System Reset Decision (00:30 - 00:45)

- User authorized full reinstall of OpenHuman
- Uninstall completed via “Add or remove programs”
- Data wipe: Removed .openhuman, com.openhuman.app, registry entries
- Prepared for clean reinstall with UI-first approach

TECHNICAL ARTIFACTS

Configuration Files Analyzed

- config.toml (OpenHuman user profile, 18KB, TOML format)
- active_user.toml (user ID pointer)
- openhuman.2026-07-09.log (application logs, JSON-RPC format)
- config.ini (AgenticSeek, INI format)
- agent.py (AgenticSeek base agent class, Python)

Commands Executed

- taskkill /f /im OpenHuman.exe (process termination)
- rmdir /s /q (directory removal)
- notepad config.toml (manual config editing)
- git clone https://github.com/Fosowl/agenticSeek.git (repo cloning)
- python patch_governance.py (automated code patch)
- Get-WmiObject Win32_VideoController (GPU detection)
- ollama list (local model enumeration)

Models in Ollama Inventory

- dhu-ornith-builder:latest
- ornith:9b
- ornith:35b

- gemma4:latest
- llama3:latest
- dolphin-llama3:latest
- mistral:latest
- qwen2.5-coder:14b
- (and others – 30+ models confirmed)

OBSERVATIONS

1. OpenHuman v0.58.7 uses a hardcoded MVP allowlist for chat models, silently falling back to cloud models when custom models are selected. This is architectural, not a bug.
2. The Electron window state system is fragile – 8 zombie processes accumulated from failed launches, requiring manual taskkill intervention.
3. AgenticSeek's Python architecture allows prompt-level governance injection, making it superior for framework implementation.
4. The Gen-Cultivate framework produced measurable behavioral change when ingested into AI reasoning – responses became bounded, precise, and authority-preserving.
5. The user's hardware (RTX 4070 Ti) is capable of running 35B parameter models locally with GPU acceleration.

NEXT STEPS

1. Complete OpenHuman reinstall (UI-first approach)
2. Verify model selection via interface (not config file)
3. If blocked, pivot to AgenticSeek as primary platform
4. Document governance injection results
5. Prepare public-safe portfolio materials

Log Authority: Ramon E. Ramirez **Session Date:** 2026-07-09/10 **Classification:** Career Documentation / Technical Evidence